

DECISION LOG

Every build choice, documented as it was made

ShopMind | Track 1: AI Shopping Agent | Kasparro Agentic Commerce Hackathon | April 2026

What this log is

A running record of every significant decision made during the ShopMind build. Format: what we considered, what we chose, and why. Written during the build, not retrospectively. This is the hardest signal to fake and the most useful one to read.

Decision Log — ShopMind Build

ID	Decision point	What we chose	Why — and the tradeoff accepted
D01	Shopify API: Storefront vs Admin	Admin API (GraphQL)	Storefront API lacks access to custom metafields — which are the entire intelligence layer. Admin API requires server-side calls (security tradeoff), but the metafield access is irreplaceable. Tradeoff accepted.
D02	Track selection: which of 5 tracks	Track 1: AI Shopping Agent	Track 1 has the clearest evaluation criteria and the most direct link to Kasparro's stated mission. Track 5 (AI Representation Optimizer) was interesting but requires real merchant data we don't have. Track 1 lets us build the full decision pipeline on synthetic data.
D03	Category: gifting vs skincare vs fashion	Gifting	Gifting has the highest emotional complexity — the hardest test of L1 (emotional intent parsing). If the system works for gifting, it works for simpler categories. Also: gifting has the most obvious return rate problem (regret), making

			the business case concrete.
D04	How many products in the catalog	~40 synthetic products	Enough variety to make L2 narrowing meaningful (8-12 products per occasion type). Small enough to hand-tag metafields properly. Large enough to demonstrate that ShopMind narrows, not just matches.
D05	One question per turn vs batch questions	One question per turn, hard constraint	Tested both approaches in prompt design. Multi-question turns felt interrogative in conversation testing. The single question rule is harder to implement (requires the LLM to prioritise which dimension matters most) but produces a dramatically better buyer experience. Tradeoff: more turns needed.
D06	Show confidence score to buyer vs hide it	Show it — framed as 'match strength'	Hiding it felt dishonest and removed a natural escalation trigger (below 75%: ask one more question). Framing as 'match strength' rather than 'system confidence' tested well — it's about the product fit, not the AI's uncertainty. Risk: might reduce trust if score is low.
D07	Mood canvas: text tiles vs image tiles	Text tiles (words like 'warm', 'minimal', 'bold')	Image tiles require asset creation and introduce inconsistent aesthetic interpretation. Text tiles map cleanly to emotion_tags in the metafield schema. Lower implementation cost, more direct mapping, better for V1 validation.
D08	LLM for cart operations vs deterministic	Deterministic — no exceptions	LLMs hallucinate. Cart operations touch real inventory and real money. The LLM reasons about what to recommend; deterministic code executes every Shopify mutation. This boundary is non-negotiable.

D09	L10 persistence: session vs cross-session	Session-scoped by default, opt-in for persistence	DPDP (India's data protection law, 2023) requires consent for personal data storage. Recipient preferences are personal data. Session-scoped by default is compliant by design. Opt-in persistence is DPDP-compliant and gives the merchant a legitimate LTV signal.
D10	Social proof: aggregate ratings vs persona-matched	Persona-matched review metafields	Aggregate ratings (4.2 stars, 847 reviews) tell the buyer nothing about whether this product is right for someone like them. A review written by 'a daughter buying for a picky mom' is 10x more persuasive. Requires the review_persona metafield — more setup, dramatically better signal.
D11	Upsell timing: before vs after recommendation	After recommendation, before cart add	Upsell before the main recommendation feels pushy. Upsell after cart add is too late (decision is closed). The window is: after recommendation + confidence is high + before cart. L9 fires exactly here.
D12	Failure fallback: general apology vs specific action	Specific action for every failure type	A general 'something went wrong' message is useless. Each failure has a defined response: API down = graceful pause, malformed JSON = retry + fallback question, invalid product_id = re-prompt with catalog. The buyer never sees a broken state.
D13	Documentation format: markdown vs Word	Both — MD for repo, Word/PDF for submission	DECISION_LOG.md lives in the repo (git-trackable, diffable, readable on GitHub). Product and Technical Documents are PDF (submission format). The content is identical; the format serves the audience.
D14	Demo approach: live	Recorded walkthrough with	Live demos fail at the

	demo vs recorded walkthrough	narration	worst moments. A 3-5 min recorded walkthrough with narration demonstrates every logic in sequence, can be edited to show the best path, and is replayable by judges. Supplemented with repo code for verification.
--	------------------------------	-----------	--

Log status:

14 decisions documented. Log will be updated during the remaining build period. Final version committed with submission. All decisions written at the time they were made — not reconstructed after the fact.

A decision log is a thinking record, not a justification record.

We wrote these as we built — not after we knew what worked.

That's the difference between product thinking and product theatre.